



# **Using MySQL 5.6 (and above) with Policy Manager 8.0 and Community Manager 8.0**

## Using MySQL 5.6 (and above) with Policy Manager 8.0 and Community Manager 8.0

using\_mysql\_56\_pm80\_cm80

January 2016

### Copyright

Copyright © 2016 Akana, Inc. All rights reserved.

### Trademarks

Akana, SOA Software, Policy Manager, Portfolio Manager, Repository Manager, Service Manager, Community Manager, Akana Intermediary for Microsoft and SOLA are trademarks of Akana, Inc. All other product and company names herein may be trademarks and/or registered trademarks of their registered owners.

### Akana, Inc.

Akana, Inc.

12100 Wilshire Blvd, Suite 1800

Los Angeles, CA 90025

(866) SOA-9876

[www.akana.com](http://www.akana.com)

[info@akana.com](mailto:info@akana.com)

### Disclaimer

The information provided in this document is provided “AS IS” WITHOUT ANY WARRANTIES OF ANY KIND INCLUDING WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, OR NON-INFRINGEMENT OF INTELLECTUAL PROPERTY. Akana may make changes to this document at any time without notice. All comparisons, functionalities and measures as related to similar products and services offered by other vendors are based on Akana’s internal assessment and/or publicly available information of Akana and other vendor product features, unless otherwise specifically stated. Reliance by you on these assessments / comparative assessments is to be made solely on your own discretion and at your own risk. The content of this document may be out of date, and Akana makes no commitment to update this content. This document may refer to products, programs or services that are not available in your country. Consult your local Akana business contact for information regarding the products, programs and services that may be available to you. Applicable law may not allow the exclusion of implied warranties, so the above exclusion may not apply to you.

## Introduction

If you are using MySQL 5.6 or higher with Policy Manager 8.0 and/or Policy Manager 8.0 and Community Manager 8.0, the following MySQL property settings must be configured as outlined in the following sections.

## MySQL Properties

### innodb\_file\_per\_table

When `innodb_file_per_table` is enabled (the default in MySQL 5.6.6 and higher), InnoDB stores the data and indexes for each newly created table in a separate `.ibd` file, rather than in the system tablespace. The storage for these InnoDB tables is reclaimed when the tables are dropped or truncated.

`innodb_file_per_table` is dynamic and can be set ON or OFF using `SET GLOBAL`. You can also set this parameter in the MySQL configuration file (`my.cnf` or `my.ini`) but this requires shutting down and restarting the server.

Akana requires this setting enabled so tables can use DYNAMIC row format.

```
set global innodb_file_per_table=ON;
```

### Advantages

- Tables created in file-per-table tablespaces use the Barracuda file format. The Barracuda file format enables features such as compressed and dynamic row formats. Tables created in the system tablespace cannot use these features.
- You can reclaim disk space when truncating or dropping a table stored in a file-per-table tablespace. Truncating or dropping tables stored in the system tablespace creates free space internally in the system tablespace data files (ibdata files) which can only be used for new InnoDB data.
- The `TRUNCATE TABLE` operation is faster when run on tables stored in file-per-table tablespaces.
- You can run `OPTIMIZE TABLE` to compact or recreate a file-per-table tablespace. When you run an `OPTIMIZE TABLE`, InnoDB creates a new `.ibd` file with a temporary name, using only the space required to store actual data. When the optimization is complete, InnoDB removes the old `.ibd` file and replaces it with the new one. If the previous `.ibd` file grew significantly but the actual data only accounted for a portion of its size, running `OPTIMIZE TABLE` can reclaim the unused space.
- File-per-table tablespaces may improve chances for a successful recovery and save time when a corruption occurs, when a server cannot be restarted, or when backup and binary logs are unavailable.
- You can back up or restore individual tables quickly using the MySQL Enterprise Backup product, without interrupting the use of other InnoDB tables. This is beneficial if you have tables that require backup less frequently or on a different backup schedule.
- File-per-table tablespaces are convenient for per-table status reporting when copying or backing up tables.
- You can monitor table size at a file system level, without accessing MySQL.

- The system tablespace stores the data dictionary and undo logs, and has a 64TB size limit. By comparison, each file-per-table tablespace has a 64TB size limit, which provides you with room for growth.

## Disadvantages

- With file-per-table tablespaces, each table may have unused space, which can only be utilized by rows of the same table. This could lead to wasted space if not properly managed.
- `fsync` operations must run on each open table rather than on a single file. Because there is a separate `fsync` operation for each file, write operations on multiple tables cannot be combined into a single I/O operation. This may require InnoDB to perform a higher total number of `fsync` operations.
- **mysql** must keep one open file handle per table, which may impact performance if you have numerous tables in file-per-table tablespaces.
- If many tables are growing there is potential for more fragmentation which can impede `DROP TABLE` and table scan performance. However, when fragmentation is managed, having files in their own tablespace can improve performance.

## `innodb_large_prefix`

By default, an index key for a single-column index can be up to 767 bytes. When this option is enabled, index key prefixes longer than 767 bytes (up to 3072 bytes) are allowed for InnoDB tables that use the `DYNAMIC` and `COMPRESSED` row formats.

Akana requires this setting enabled since it uses index keys that are greater than 767 bytes in length.

```
set global innodb_large_prefix=ON;
```

## `innodb_file_format`

The file format to use for new InnoDB tables. Currently, Antelope and Barracuda are supported. This setting only applies to tables that have their own file-per-table tablespace, so for it to have an effect, `innodb_file_per_table` must be enabled. The Barracuda file format is required to use Compressed or Dynamic row formats and associated features such as compression, off-page storage for large variable-length columns, and large index key prefixes.

Akana requires that the file format be set to Barracuda so tables can be created with `DYNAMIC` row format.

```
set global innodb_file_format=barracuda;
```

## Advantages

- The `DYNAMIC` row format maintains the efficiency of storing long data values such as BLOB and TEXT.

## **max\_allowed\_packet**

The maximum size of one packet or any generated/intermediate string, or any parameter sent by the application to the server. The default is 4MB.

You must increase this value if you are using large BLOB columns or long strings. It should be as big as the largest BLOB you want to use. The protocol limit for `max_allowed_packet` is 1GB. The value should be a multiple of 1024; non-multiples are rounded down to the nearest multiple.

Akana requires the packet size set to 1073741824 (1GB) since BLOB and TEXT columns are used.

```
set global max_allowed_packet=1073741824;
```

## **innodb\_flush\_log\_at\_trx\_commit**

Controls the balance between strict ACID compliance for commit operations, and higher performance that is possible when commit-related I/O operations are rearranged and done in batches. You can achieve better performance by changing the default value, but then you can lose up to a second of transactions in a crash.

- The default value of 1 is required for full ACID compliance. With this value, the contents of the InnoDB log buffer are written out to the log file at each transaction commit and the log file is flushed to disk.
- With a value of 0, the contents of the InnoDB log buffer are written to the log file approximately once per second and the log file is flushed to disk. No writes from the log buffer to the log file are performed at transaction commit. Once-per-second flushing is not 100% guaranteed to happen every second, due to process scheduling issues. Because the flush to disk operation only occurs approximately once per second, you can lose up to a second of transactions with any **mysqld** process crash.
- With a value of 2, the contents of the InnoDB log buffer are written to the log file after each transaction commit and the log file is flushed to disk approximately once per second. Once-per-second flushing is not 100% guaranteed to happen every second, due to process scheduling issues. Because the flush to disk operation only occurs approximately once per second, you can lose up to a second of transactions in an operating system crash or a power outage.

Akana does not require this setting changed from default value '1'. However if you are seriously looking at improving the performance, the value can be changed to '0' or '2'.

```
set global innodb_flush_log_at_trx_commit=0;
```